

# Report on Linear Regression Model Training and Pipeline

## 1. Objective

The primary goal of this project was to build, evaluate, and compare the performance of a Linear Regression model using two different approaches:

1. Using the **LinearRegression** class from the scikit-learn package.
2. Implementing the Linear Regression model from scratch using Gradient Descent.

The comparison was based on the Mean Squared Error (MSE) and R-squared ( $R^2$ ) metrics.

## 2. Dataset

The dataset used is the "Bike Sharing Dataset," which records the hourly count of rental bikes in a city along with various weather and seasonal information.

### Key Features:

- **Categorical:** `season`, `weathersit`, `mnth`, `yr`, `hr`, `holiday`, `weekday`, `workingday`, and `day_night` (derived from the `hr` feature).
- **Numerical:** `temp`, `hum`, `windspeed`, `temp_hum`, and `temp_windspeed` (interaction terms).

### Target Variable:

- **cnt:** The total count of bike rentals during each hour.

## 3. Data Preprocessing

The preprocessing involved handling both categorical and numerical features:

### Numerical Features:

- **Imputation:** Missing values were imputed using the mean.
- **Scaling:** All numerical features were scaled using the **MinMaxScaler** to bring them into the same range.

## Categorical Features:

- **Imputation:** Missing values were imputed using the most frequent value.
- **Encoding:** Categorical variables were encoded using `TargetEncoder`, which encodes categories with the mean of the target variable for each category.

## 4. Model Training

### a. Using the `LinearRegression` Package

The Linear Regression model was trained using the `LinearRegression` class from scikit-learn. This model was fitted to the preprocessed training data, and predictions were made on the test data.

### b. Implementing Linear Regression from Scratch

The model was implemented using the following steps:

1. **Adding a Bias Term:** An intercept term was added to the features.
2. **Normalization:** Features were normalized using `StandardScaler`.
3. **Gradient Descent:**
  - Weights were initialized to zero and iteratively updated using the gradient descent algorithm to minimize the MSE.
4. **Prediction:** Predictions were made on the test set using the learned weights.

## Evaluation Metrics

- **Mean Squared Error (MSE):** Measures the average squared difference between actual and predicted values.
- **R-squared ( $R^2$ ):** Indicates the proportion of the variance in the target variable that is predictable from the features.

## Results Comparison

Metric	Using Package	Scratch Implementation
Mean Squared Error (MSE)	14,974.13	49,549.98
R-squared ( $R^2$ )	0.527	-0.565

### Observations:

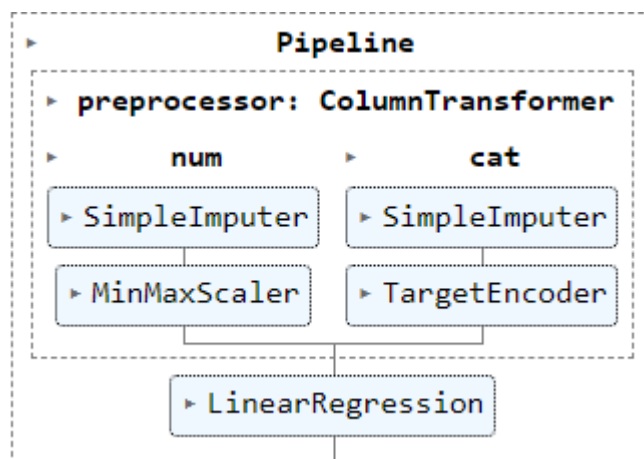
- The model trained using the `LinearRegression` package performed significantly better, with a lower MSE and higher  $R^2$ , indicating a better fit.
- The scratch implementation had a higher MSE and a negative  $R^2$ , suggesting that the model did not perform well, likely due to challenges in optimizing the gradient descent algorithm.

## 5. Pipeline Visualization

Below is the visualization of the ML pipeline that combines both preprocessing and model training steps:

### Explanation:

- The `preprocessor` step combines numerical and categorical preprocessing.
- The `model` step fits a Linear Regression model to the preprocessed data.



## 6. Conclusion

This project successfully demonstrated the process of building and evaluating a Linear Regression model using two different approaches. The scikit-learn implementation outperformed the manual implementation, highlighting the importance of using well-optimized libraries for machine learning tasks.